

ChatGPT history

You said:

Best gpu to get the best flux at home image manipulation. Along with model sizes/names

ChatGPT said:

If you want the “best GPU” for home use with what you refer to as “flux / image manipulation” (I assume you mean modern AI-based image generation / editing, e.g. using models like FLUX.1 or similar diffusion / generative-AI models), then you want a GPU with *lots of VRAM, modern tensor-cores*, and good memory bandwidth — because these determine how big models you can load and how fast your generation will run. [Propel RC+2Tom's Hardware+2](#)

Here are GPUs widely recommended today (2025), across different budget/performance levels:

✓ Top GPUs for AI Image Generation at Home (2025)

NVIDIA GeForce RTX 3090 24 GB GDDR6X

€1,447.60

•

Router-switch.com

PNY NVIDIA RTX 5000 ADA Graphics Card

€4,669.69

•

Cdiscount + others

NVIDIA RTX 4000 SFF Ada 20 GB Graphics Card

€1,458.85

•

Senetic.fr + others

PNY Nvidia 2000 ADA 16 GB Graphics Card

€822.34

•

PcComponentes FR + others

NVIDIA L4 24 GB GDDR6 GPU

€2,744.06

•

Senetic.fr + others

HP NVIDIA Quadro M5000 GPU Module

€690.20

•

it-market.com

GIGABYTE AMD Radeon RX 9060 XT Graphics Card

€324.99

•

AiO

Acer Nitro Radeon RX 9060 XT OC Graphics Card

€429.90

•

Acer Store France + others

High-end / Future-proof (Large models, high-res, heavy workloads)

- **NVIDIA GeForce RTX 3090 24 GB GDDR6X** — 24 GB VRAM still makes this a classic for large models (SDXL, FLUX, big batch sizes, high-res output). Good trade-off if you want top-level performance without going full workstation.
- **PNY NVIDIA RTX 5000 ADA Graphics Card** — More of a “prosumer / workstation” card: likely massive VRAM / compute headroom, useful if you plan heavy experimentation, many generated images in batch, or maybe some training/fine-tuning.

Balanced Mid-to-High (Good for most use-cases; efficient & versatile)

- **NVIDIA RTX 4000 SFF Ada 20 GB Graphics Card** — 20 GB VRAM and good tensor/core architecture give plenty of room for modern diffusion models, even larger ones or multiple models loaded simultaneously.
- **PNY Nvidia 2000 ADA 16 GB Graphics Card** — 16 GB VRAM is often a “sweet spot” for many current-generation diffusion models if you don’t push extreme high-res or batch sizes. Good price-to-performance ratio for creative AI workflows.

Professional / Inference-oriented (For heavier workloads / multipurpose)

- **NVIDIA L4 24 GB GDDR6 GPU** — More enterprise-class / compute-oriented GPU; if you are serious about heavy image generation or long sessions, this gives VRAM and stability that help avoid bottlenecks.
- **HP NVIDIA Quadro M5000 GPU Module** — Older workstation-class GPU; could work if you find good second-hand deals, though not state-of-the-art for 2025.

Budget / Entry-level (Good for experimentation, smaller models)

- **GIGABYTE AMD Radeon RX 9060 XT Graphics Card** — Lower-cost GPU, okay for lightweight or quantized versions of models, smaller images or lower resolutions; not ideal for heavy workloads, but fine to start experimenting.
 - **Acer Nitro Radeon RX 9060 XT OC Graphics Card** — Similar to above; acceptable for entry-level and “just try things out” workflows.
-

What “Model Sizes / Names” you can reasonably run with different GPUs

- With 24 GB+ VRAM (like RTX 3090 or workstation GPUs): you can comfortably run large models (e.g. high-capacity versions of FLUX, high-res output, larger batch sizes, maybe even moderate fine-tuning). [poolcompute.com+2UMA Technology+2](#)
 - With 16 – 20 GB VRAM: good for most modern diffusion models — including mid-size and many “SDXL/FLUX-style” ones — especially if you stick to single-image generation or moderate-resolution outputs. [iRender Cloud Rendering Service+1](#)
 - With ~8–12 GB VRAM (budget GPUs): you’re limited to smaller or quantized models, lower resolutions, maybe slower generation, and likely no large batch processing. Works well for experimentation. [XDA Developers+1](#)
-

My Recommendation Based on Your Use

Since you’re in Paris and presumably building a workstation for “home image manipulation / AI generation”, if you want **top performance and flexibility**, go for an RTX 3090 or a modern workstation-class 24 GB+ card.

If you want a **balanced, efficient setup** that’s powerful but more reasonable in price/power use, a 16–20 GB card (like the PNY 2000 ADA or RTX 4000-class) is often the most practical “sweet spot.”

If you are just **starting out / experimenting**, a budget GPU with 8–12 GB is acceptable — but expect limitations with big models or high resolution.

If you want — I can **pull up 5–10 of the best value-for-money GPUs in 2025** for AI image generation (i.e. best performance per euro) **targeted for users in Europe (availability / price-to-performance optimized)**.

Do you want me to build that list for you now?

You said:

I have a rtx3060 12 gig. A second machine with two gtx1080 ti What is the next step up? I use flex/comfy

ChatGPT said:

Given that you already have a NVIDIA RTX 3060 12 GB and a second machine with two GTX 1080 Ti — you're solid for entry- and mid-level AI image generation (ComfyUI / FLUX etc.).

Aiarty+2Best GPUs for AI+2

If you want to upgrade now to get significantly faster generation, support bigger models, higher-res output, or smoother batch workflow, then stepping up to a “next-tier” GPU makes sense. Here are some good upgrade options (sorted by performance / cost / headroom trade-off):

Good Upgrade GPUs to Consider

GIGABYTE RTX 4090 GAMING OC 24GB GDDR6X

€2,179.90

•

instinctgaming.gg + others

NVIDIA GeForce RTX 3090 24 GB GDDR6X

€1,447.60

•

Router-switch.com

MSI GeForce RTX 4070 Ti VENTUS 3X E1 12G OC

€708.20

•

PcComponentes FR

NVIDIA GeForce RTX 4070 12GB GDDR6X

€600.03

•

AliExpress + others

NVIDIA GeForce RTX 3080 Ti 12GB GDDR6X

€816.00

•

StockX + others

GIGABYTE GeForce RTX 3080 10GB GDDR6X

€395.90

•

Computer Store Berlin OHG + others

GeForce RTX 4070 SUPER VERTO 12GB GDDR6X

€521.96

- pskmegastore + others

RTX 4070 Ti ASUS ROG STRIX OC 12GB

€728.73

- PcComponentes FR + others

Highlights

- **GIGABYTE RTX 4090 GAMING OC 24GB GDDR6X** — top-tier consumer GPU today. With 24 GB VRAM and a powerful tensor core setup, it gives you **maximum headroom**: fast image generation, ability to load large models or multiple models at once, smooth high-res output (e.g. 2048×2048 or more), bigger batches, and future-proofing for upcoming heavy models. Great if you want “do-everything” performance at home.
- **NVIDIA GeForce RTX 3090 24 GB GDDR6X** — slightly older but still very strong, with 24 GB VRAM. If you can find one at a good price (possibly used or discounted), it remains a **workhorse for large-model inference** and heavy AI workloads.
- **MSI GeForce RTX 4070 Ti VENTUS 3X E1 12G OC** — a more balanced “value-performance” option. Good step up from a 3060: newer architecture, increased compute/Tensor-core efficiency, and decent VRAM for many models. Good for mid-size to heavy workloads without going into flagship price territory.
- **NVIDIA GeForce RTX 4070 12GB GDDR6X** — a moderate-budget modern upgrade. If you don’t need the absolute max VRAM, this gives you architectural improvements over 3060 and GTX 1080 Ti. It’s a solid choice if you want smoother performance with recent models while staying in a modest budget.
- **NVIDIA GeForce RTX 3080 Ti 12GB GDDR6X** — older high-end card; decent compute power and probably faster than your 3060/1080 Ti in many tasks. Worth considering if you find it at a good price.
- **GIGABYTE GeForce RTX 3080 10GB GDDR6X** — less VRAM than ideal for big AI models, but if your workflows stay moderate (smaller models, lower-res images), it can still provide a tangible speed boost over 3060/1080 Ti.
- **GeForce RTX 4070 SUPER VERTO 12GB GDDR6X** — another modern mid-range GPU with good performance per watt and relatively up-to-date architecture; a fine pick if budget vs performance matters.
- **RTX 4070 Ti ASUS ROG STRIX OC 12GB** — comparable to the other 4070-class cards; decent “sweet-spot” choice for users upgrading from 3060 without aiming for flagship-level spend.



What You Gain vs Your Current Setup

- More **VRAM (16–24 GB vs your 12 GB / 12 GB)** → you'll be able to run larger models, high-res outputs, and batch generation more reliably and without memory errors. [Aiarty+2TechBloat+2](#)
 - **Faster generation times:** newer architecture + more tensor cores → less waiting per image, smoother workflow, especially with heavier/modern models. [Best GPUs for AI+1](#)
 - **Better future-proofing:** newer models are trending towards more VRAM use (larger diffusion models, bigger context, more complex pipelines). Upgrading now means fewer painful upgrades later.
-



My Recommendation (Based on Your Use)

Since you already run FLUX/Comfy-based workflows:

- If you want **top performance and long-term headroom**, go for **RTX 4090** — ideal if you want high-res, batch processing, or plan to dive into heavy fine-tuning or video-generation.
 - If you want a **good balance** between cost, performance, and VRAM — a **4070 Ti / 4070** is a very reasonable and efficient next step.
 - If budget is moderate and you just want a noticeable boost over 3060/1080 Ti — **3080 / 3080 Ti (if priced well)** or **4070** are pragmatic picks.
-

If you like — I can run a **2025 price vs performance comparison table** (for the European market) — showing *cost per VRAM per FLOP per watt* for ~5-10 of the most relevant GPUs for AI/image-generation.

Would you prefer I build that table for you now (given your Paris location)?

You said:

My goal is photorealistic. Not speed

ChatGPT said:

Good clarification. If your goal is **photorealistic image generation** rather than raw speed, then VRAM capacity and model compatibility matter more than “iterations per second.” With that in mind — given your current hardware (RTX 3060, two GTX 1080 Ti) — here are what I’d consider your “**next-step up**” **GPUs** for maximum flexibility, headroom, and ability to push **high-quality / high-resolution / complex models** (e.g. heavy variants of ComfyUI + FLUX or big diffusion models).

🎯 Best GPUs for Photorealistic Image Generation (Upgrade Picks)

NVIDIA GeForce RTX 4090 24 GB GDDR6X

€876.53

•

AliExpress

ASUS ROG Strix NVIDIA GeForce RTX 3090 24 GB GDDR6X

€34.90

•

ColdZero + others

NVIDIA GeForce RTX 4080 16 GB GDDR6X

€1,469.99

•

Amazon.fr - Seller + others

MSI GeForce RTX 4070 Ti VENTUS 3X 12 GB GDDR6X

€708.20

•

PcComponentes FR

ASUS TUF NVIDIA GeForce RTX 3080 Ti 12 GB GDDR6X

€790.96

•

PcComponentes FR + others

NVIDIA GeForce RTX 3080 Ti 12 GB GDDR6X (StockX)

€816.00

•

StockX + others

NVIDIA GeForce RTX 3080 10 GB GDDR6X Founders Edition

€579.45

•

AliExpress + others

GIGABYTE NVIDIA GeForce RTX 3090 24 GB GDDR6X (OC)

€1,413.99

•

Rakuten - Informatique - Occasion + others

🏅 Elite / “Do-Anything” Option

- **NVIDIA GeForce RTX 4090 24 GB GDDR6X** — 24 GB VRAM and modern architecture give you maximum headroom. Ideal if you want to produce large, high-resolution, photorealistic images, work with complex pipelines (LoRA, ControlNet, upscaling, etc.), or keep your setup future-proof.

Workhorse “Pro-Level but More Affordable” Option

- **ASUS ROG Strix NVIDIA GeForce RTX 3090 24 GB GDDR6X** — still 24 GB VRAM, enough to run heavy models and high-res outputs. Performance slightly behind 4090, but often a better value if found at a good price, especially for strictly image generation (not cutting-edge video).

Balanced Newer-Gen Option (Good trade-off speed/VRAM/price)

- **NVIDIA GeForce RTX 4080 16 GB GDDR6X** — 16 GB VRAM is often enough for high-res photorealistic images (especially if you manage memory / VRAM usage carefully), and newer Ada architecture brings efficiency and better tensor-core support for diffusion-based workflows.
- **MSI GeForce RTX 4070 Ti VENTUS 3X 12 GB GDDR6X** — a more modest upgrade relative to your 3060. Not the top VRAM, but a reasonable middle-ground if you expect moderately heavy photorealistic outputs without going full flagship.

Older High-End (Value / Budget-conscious) Options

- **ASUS TUF NVIDIA GeForce RTX 3080 Ti 12 GB GDDR6X or NVIDIA GeForce RTX 3080 Ti 12 GB GDDR6X (StockX)** — if you find them at good prices, they can still handle many photorealistic workloads. VRAM is decent, though more limited than 24 GB cards, so you might need to optimize (lower res, fewer control layers).
- **NVIDIA GeForce RTX 3080 10 GB GDDR6X Founders Edition** — not ideal for top-tier photorealism (10 GB VRAM becomes a limiting factor for large models or high-res). Only recommended if price is very low and you are comfortable compromising on resolution / complexity.
- **GIGABYTE NVIDIA GeForce RTX 3090 24 GB GDDR6X (OC)** — similar to the other 3090; a good “workhorse 24 GB” with modest overclocking headroom.

Why These Are Worth Upgrading — Even for Photorealism (Not Just Speed)

- **VRAM matters for photorealism:** According to GPU-for-AI-image-gen guides, 24 GB VRAM GPUs give you “full power” — enough headroom for high-resolution, multiple condition networks (e.g. ControlNet + LoRA + upscaling), batch runs, or complex model stacks. [HOSTKEY — premium web services provider+2NunarIQ+2](#)

- **VRAM + architecture = flexibility:** Newer-gen GPUs (Ada / Ampere) support FP16/BF16/FP8 quantization, tensor cores, and more efficient memory bandwidth — enabling heavy models to run smoothly and producing high-quality outputs without crashing on out-of-memory. [ComfyUI Wiki+1](#)
 - **Future-proofing:** AI/image generation models continue evolving: higher-res outputs, more conditioning layers, multiples models in one generation — 24 GB (or more) VRAM gives breathing room for future workloads. [HOSTKEY — premium web services provider+1](#)
-

🎯 My Recommendation for Your Use-Case (Photorealistic + Flex/Comfy)

Since you're not chasing speed but photorealistic quality — I'd aim for **24 GB-class GPU as next step**, to maximize flexibility and reduce limitations when using heavy photorealistic models (especially if you experiment with higher resolution or control networks).

If budget is a concern but you still want a meaningful boost over your current hardware, a **16 GB (or slightly lower) modern GPU** can still work — though you'll need to manage VRAM carefully (lower res, simpler pipelines).

If you want — I can **project 2–3 realistic “upgrade paths”** (low-cost, mid-budget, high-end) — with approximate cost in € (Europe/France 2025) and expected capability (res, batch size, model types).

You said:

What should I look for Black Friday cyber Monday. 3090s?

ChatGPT said:

Good question. Because your goal is **photorealistic image generation** (not necessarily fastest generation) — Black Friday / Cyber Monday can be a good time to look for an upgrade GPU. But there are a few things you should watch out for, especially with high-VRAM cards like ZOTAC GAMING GeForce RTX 3090 Trinity OC 24GB GDDR6X or other 24 GB cards.

Here's **what you should look for and how to approach Black Friday for your use-case** 🤝

✓ What to check during Black Friday / Cyber Monday

- **True discount vs. price inflation** — Many “Black Friday deals” simply revert GPUs to MSRP (or slightly below), so a 5-10% off isn’t always worth it unless the card is exactly what you need. [Digital Trends+1](#)
 - **VRAM / memory headroom** — Since you prioritize photorealism and likely larger models / high-res outputs, prefer cards with **≥ 24 GB VRAM**. That means older high-end cards (like the 3090) or the top flagship GPUs.
 - **Architecture & support for AI workloads** — Tensor cores, FP16/BF16 support, and good driver/optimization support can matter more than raw gaming throughput if you mainly run diffusion / generative-AI pipelines.
 - **Realistic price floor (or below MSRP)** — For a card like the 3090 or similar, the “deal” is only worth it if the price is noticeably below what it usually sells for (or near the historical MSRP). Otherwise you may not save much. [TechSpot+1](#)
 - **Stock / availability & PSU/compatibility** — Make sure the card fits in your case, your PSU can supply enough power, and that supply/shipping logistics won’t eat up the gains. [Powers Of 10+1](#)
 - **Avoid hype-driven “discounts” on new mid-range cards when you actually need VRAM headroom** — Many deals focus on mid-range GPUs (good for gaming) — but for model-heavy workloads you want heavy-VRAM cards. [Windows Central+1](#)
-

🎯 Should you hunt for 3090s or other cards this Black Friday?

Yes — **if** you can find a 24 GB-class GPU at a *good* price. For photorealistic AI-generation, older high-end GPUs (like 3090) can still be among the smartest purchases, because their VRAM and memory bus/width give you headroom many new mid-range GPUs lack.

For example:

- **ZOTAC 3090 24 GB GDDR6X** — 24 GB VRAM gives you flexibility for large or multiple models, high-res outputs, or heavy pipelines.
- A high-end 24 GB card from another brand or OEM — as long as it’s not overpriced — remains a strong bet, especially if your workflow includes large photorealistic outputs, ControlNets, LoRAs, or chaining models.

However — don’t treat all Black Friday GPU deals as equal. If the only “discounts” are on lower-VRAM mid-range GPUs (12–16 GB), they may be worthless for your use case — unless you accept limitations in resolution, model size, or batch complexity.

Also, some sources argue that GPU discounts during Cyber Week are often modest, or just restore MSRP rather than offer deep savings. [Digital Trends+1](#)



What you should do right now (before Black Friday ends)

- Monitor trusted retailers (especially those shipping to France/Europe) for **24 GB+ GPUs** — old-gen but VRAM-heavy cards.
 - Compare historical pricing to sale price — check price trackers or prior listings to know if the deal is real.
 - Confirm your PSU, case, and power-requirements before purchasing.
 - If you find a good 3090 (or equivalent) from a **reputable seller at a realistic price** — go for it (since VRAM headroom is rare in newer mid-range cards).
-

If you like, I can **scan a dozen European retailers** (France / EU) *right now* and **pull up any 24 GB+ GPUs (like 3090, 4090, etc.)** that are currently discounted — with approximate price in €. Do you want me to build that list for you?

You said:

Give me a hit list for Amazon. Com. Us.

ChatGPT said:

Good idea — a “hit list” (on Amazon.com / US) focused on GPUs that make sense for **photorealistic image generation** (i.e. high VRAM or good compute for diffusion / ML-based image workloads). Below are several GPUs I’d watch for in Black Friday / Cyber Monday sales — ordered roughly from “best overall for photorealism” down to “budget / fallback” options.

✓ GPUs to Watch on Amazon.com (for Photorealistic / AI Image Workflows)

NVIDIA GeForce RTX 3090 24 GB GDDR6X

€1,447.60

•

Router-switch.com + others

NVIDIA GeForce RTX 4080 16 GB GDDR6X

€1,469.99

•

Amazon.fr - Seller + others

GIGABYTE GeForce RTX 3080 Ti 12 GB GDDR6X

€781.44

•

PcComponentes FR + others

PNY NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X

€1,070.42

•

Carrefour + others

ASUS TUF NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X

€1,183.08

•

OnBuy.com + others

GeForce RTX 4070 SUPER 12 GB GDDR6X

€521.96

•

pskmegastore + others

ASUS TUF GeForce RTX 3080 10 GB GDDR6X

€462.40

•

PcComponentes FR + others

GIGABYTE GeForce RTX 3080 10 GB GDDR6X Gaming OC

€350.00

•

Leboncoin + others

↑ Premium / “Do-Anything” Options

- **NVIDIA GeForce RTX 3090 24 GB GDDR6X** — With 24 GB VRAM, this remains one of the best “all-around” cards for heavy diffusion / generative-AI workflows, high-res outputs, or stacking multiple models/conditioners.
- **NVIDIA GeForce RTX 4080 16 GB GDDR6X** — Newer-generation architecture (from Ampere → Ada or similar) gives efficiency and modern tensor-core support. 16 GB VRAM is lower than 3090 but might suffice if you optimize memory usage — good if you want newer architecture over raw VRAM.

⚖️ Balanced / Versatile Picks (Good Performance per \$)

- **GIGABYTE GeForce RTX 3080 Ti 12 GB GDDR6X** — A solid performer if you find a good deal; still capable of realistic output, though VRAM is more limited compared to 3090.
- **PNY NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X** — Newer-gen, efficient, and decent for mid-range to advanced workloads without going flagship.
- **ASUS TUF NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X** — Similar to the above 4070 Ti, useful when you want a balance between price, power consumption, heat, and performance.

- **GeForce RTX 4070 SUPER 12 GB GDDR6X** — A more budget-conscious modern GPU; may be worthwhile if you mostly run moderate-size models or are okay optimizing VRAM usage.

Budget-oriented / Entry High-end (If Deals Are Good)

- **ASUS TUF GeForce RTX 3080 10 GB GDDR6X** — Less ideal for top-end photorealism (10 GB VRAM is limiting), but worth considering if price is much lower and you accept constraints on resolution/model size.
 - **GIGABYTE GeForce RTX 3080 10 GB GDDR6X Gaming OC** — Similar trade-offs as above: decent compute boost over older cards, but VRAM is modest for heavy AI workflows.
-

What to Watch Out for When Buying on Amazon

- **Seller legitimacy:** There have been reports of “fake/grafted” GPUs — e.g. GPUs labeled as high-end (like 4090/4080) but internally modified or mis-described. [Digital Trends](#) — Always prefer “sold by Amazon” or well-known manufacturers/resellers (Asus, Gigabyte, PNY, etc.), not random marketplace third-parties.
 - **VRAM vs architecture trade-offs:** Newer cards (like 4080, 4070 series) may have better tensor-core/AI optimization, but lower VRAM than older 24 GB-class cards; for heavy photorealistic generation you may prefer VRAM headroom.
 - **Real deal vs inflated “Black Friday” prices:** Sometimes discounts are just marketing — compare with historical price data to ensure you’re getting a genuine bargain.
 - **PSU / case compatibility:** High-end GPUs often require more power and space — ensure your system can handle the card.
-

My Advice for You (Given Your Goal: Photorealistic Output, Not Maximum Speed)

If I were you and shopping on Amazon US during Black Friday / Cyber Monday:

- I’d prioritize **24 GB-class cards** first — e.g. RTX 3090 — because VRAM headroom maximizes flexibility for high-res photorealistic jobs.
 - If price on 3090 seems excessive or stock is poor — I’d consider a newer-gen card with good tensor-core support (e.g. 4080 or 4070 Ti) as a compromise.
 - I’d avoid “bargain” 10 GB cards unless I’m ok with limiting resolution / model complexity.
-

If you want — I can scan **Amazon.com** live (US marketplace) for **all listings of 24 GB+ GPUs** (3090 / 3090 Ti / 4090 if available) — and show **price + availability** (Europe-adjusted for you if shipped to France).

Do you want me to build that list now?

You said:

Most of those make no sense. I already have a 12 and a dual 12

ChatGPT said:

You're right — given that you already own a 12 GB card (and even a dual-GPU system), some of the older “cheap mid-range” upgrades I mentioned are probably not worth it. What matters most now is **VRAM headroom, memory bandwidth / architecture**, and **real upgrade value** — especially if you want to push for photorealism, bigger/more complex models and future-proofing. [Cameraegg+2simplepod.ai+2](#)

Here's more focused thinking (concrete):

🔍 Why your current setup has real limitations

- Even though 12 GB VRAM (like on your NVIDIA RTX 3060) is okay for “basic to mid” image-generation (512×512–768×768, simpler models), it becomes a bottleneck when you:
 - want high-resolution images (1024×1024 +), especially photorealistic or heavily conditioned ones. [aihardwareprices.com+2Hyperbolic+2](#)
 - try to stack multiple models/conditioners (e.g. LoRA, ControlNet, upscaling, extra conditioning) — VRAM adds up. [Cameraegg+2simplepod.ai+2](#)
 - work with bigger or newer models (or future versions) which tend to demand more VRAM than older simple diffusion models. [bacloud.com+1](#)
- Also: two separate 12 GB-cards (e.g. your dual-GPU machine) **do not** add VRAM together for a single image generation job. VRAM doesn't stack across GPUs that way. [bacloud.com+1](#)

So: if your aim is **photorealism + flexibility + future-proofing**, a “slightly better 12 GB or 16 GB card” is often **not enough**.

✓ What kind of GPU does make sense *next* — given your current gear

If you decide to upgrade, these are the meaningful “step-up” classes that actually give you more headroom (VRAM + modern architecture), and are worth considering — rather than intermediate GPUs that only incrementally improve speed.

ASUS ROG Strix NVIDIA GeForce RTX 3090 24 GB GDDR6X
€1,295.00

•

StockX + others

Gigabyte NVIDIA GeForce RTX 3090 Ti 24 GB GDDR6X
€2,311.51

•

Cdiscount + others

MSI GeForce RTX 4090 Gaming X Trio 24 GB GDDR6X
€2,775.36

•

Router-switch.com

MSI GeForce RTX 3090 Ventus 3X 24 GB GDDR6X OC
€2,116.09

•

TecAfrica Solutions PTY + others

PNY NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X
€1,070.42

•

Carrefour + others

ASUS TUF NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X
€1,183.08

•

OnBuy.com + others

NVIDIA GeForce RTX 4070 SUPER 12 GB GDDR6X
€521.96

•

pskmegastore + others

Gigabyte GeForce RTX 4070 Ti Gaming OC 12 GB GDDR6X
€2,429.29

•

Galaxus + others

“True Upgrade — VRAM + headroom”

- **ASUS ROG Strix NVIDIA GeForce RTX 3090 24 GB GDDR6X** — 24 GB VRAM gives you the breathing room for high-res (e.g. 2K, 4K), multiple models/conditioners, heavy photorealistic workloads, or large batch renders. Still a “workhorse” in 2025 for generative-AI.

- **Gigabyte NVIDIA GeForce RTX 3090 Ti 24 GB GDDR6X** — if you find one, the 3090 Ti brings nice performance and same 24 GB headroom — good for intensive pipelines or future-proofing.
- **MSI GeForce RTX 4090 Gaming X Trio 24 GB GDDR6X** — the “top-end” (consumer) card: 24 GB VRAM + best GPU architecture, ideal if you want maximum headroom for whatever future models/practices come.
- **MSI GeForce RTX 3090 Ventus 3X 24 GB GDDR6X OC** — a slightly more affordable / OC version of 3090 if you’re buying used or discounted, still giving full VRAM headroom.

“Mid-range / modest upgrade” — probably not enough long-term

- **PNY NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X** — modern GPU & faster than your 3060, but VRAM stays 12 GB, which still limits you for high-res / heavy model stacks.
- **ASUS TUF NVIDIA GeForce RTX 4070 Ti 12 GB GDDR6X** — similar story: good efficiency and speed, but VRAM headroom is same as your existing 12 GB, so benefit is limited.
- **NVIDIA GeForce RTX 4070 SUPER 12 GB GDDR6X** and **Gigabyte GeForce RTX 4070 Ti Gaming OC 12 GB GDDR6X** — might give you slightly better speed or efficiency for “moderate” tasks, but none will dramatically expand what you *can* do.

My View: What your “Next Move” Should Be — If Photorealism & Flexibility Is the Goal

Since your aim isn’t just speed, but **serious photorealistic work + future-proofing**, I think your next upgrade should aim for **24 GB-class cards**. That means something like a 3090 (or 4090 if budget gives) — not just a “faster 12 GB card.”

Upgrading to another 12 GB/16 GB card might give smoother performance, but it won’t *unlock* new capabilities (e.g. 2048×2048+ resolution, heavy pipelines, multi-model stacks, complex conditioning, future large models). Over time, VRAM headroom becomes a real practical limit in generative-AI work. [simplepod.ai+2Cameraegg+2](#)

If you like — I can **simulate** (with approximate VRAM budgets) what a “photorealistic-ready” pipeline looks like with **12 GB, 16 GB, and 24 GB** cards — to show you concretely what you **can** vs **can’t** do in each case (resolutions, batch size, model stacks). Might help you see if your 12 GB-based setup is actually “good enough” or not. Want me to build that table for you now?

You said:

What prices are good for the 3090? Also I have a builder that gets parts. What is a used one going for. (Only interested in 24+gig vram boards.)

ChatGPT said:

Good questions — what counts as a “good price” for a NVIDIA GeForce RTX 3090 (or any 24+ GB-VRAM board) depends a bit on condition (new vs used), seller, and how much you care about warranty / wear. Based on recent market data (2025), here’s my take on what you should aim for — and what to expect if you or your builder jump on one now.

What a “Good Price” Looks Like (New or Well-Maintained Used)

- For a **used RTX 3090** in good condition (not abused by mining, with working fans, no damage, stable), a price around **€650 – €800** (or roughly **US \$700 – \$850**) is considered reasonable. [Best Value GPU+2leboncoin+2](#)
- Lower than that (say €600 / US \$650) — that could be a “very good deal,” but only if you **thoroughly inspect it** (fans, VRAM stability, booting, temps) — especially since cheap GPUs can hide issues. Many buyers on forums note that sub-€600 listings are often risky or end up being “for parts only.” [Reddit+2Reddit+2](#)
- If you see a clean and recent-generation 3090 at **~€750–€850**, that’s sort of the “sweet spot” for value vs headroom.

In short: used ~€700–€800 is very solid for a 24 GB VRAM GPU nowadays.

What You’re Likely to See on the Used / Resale Market (and Should Expect)

- On some European second-hand markets (e.g. classifieds, local resale sites), people report seeing 3090 cards at **€580–€750**, though often these go quickly and may come with caveats (shipping, testing, warranty). [leboncoin+1](#)
 - On larger marketplaces / aggregated trackers, the “average” used-price marker is around **€686** (for functional boards) — confirming that ~€650–€750 is close to what “normal” 3090s are trading for in late-2025. [Best Value GPU](#)
 - Naturally, “top-end” partner-cards (custom cooled, lightly used) fetch more — so if you find one of those for ~€800, that’s reasonable. If a listing asks ~€1200+, you’re likely paying a premium for convenience / “like new” condition — maybe less interesting unless warranty/new-box is included. [leboncoin+2Idealo+2](#)
-

⌚ My Recommendation to You (Given Your Goals & What You Already Have)

Since your goal is photorealistic image-generation and VRAM headroom matters more than “latest-gen,” I’d say:

- Try to get a used 3090 around **€700–€800** → that is where price meets value.
- Accept **€650** only if you *can check* the card thoroughly (fans, temps, stress-test) — great bargain but a little riskier.
- Avoid paying above **€900–€1000** unless the card is basically “like new,” from a reliable vendor, and comes with warranty or proof of “light use.”

Given the investment, that gives you the VRAM headroom for big models, high-res images, multi-model stacks — while still offering good value per euro.

If you want: I can **scan 5 major US + EU resale marketplaces (eBay, Reddit hardware swap, local resale, etc.)** and pull up **20–30 recent RTX 3090 listings** — with **price, condition, and seller-type**, so you get a global sense of what the “real market” looks like *today*.
Want me to build that list for you now?

You said:

Marktplatz.nl

ChatGPT said:

Good call asking about NVIDIA RTX 3090 prices — using a marketplace like Marktplaats (and related NL/EU second-hand / resale markets) is often one of the better ways to snag a 24 GB+ VRAM GPU for generative-AI / photoreal work at reasonable cost. Here’s what recent real-world listings on Marktplaats show (and what I’d consider “good prices / red-flags” based on that).

✓ What current 3090 listings look like now on Marktplaats / NL-market

From recent listing data:

Example / Description	Price (EUR) / Condition
Gigabyte GeForce RTX 3090 Gaming OC — “perfect” condition € 680 Marktplaats+1	
ZOTAC GeForce RTX 3090 Trinity 24 GB	€ 690 Marktplaats+1

Example / Description	Price (EUR) / Condition
INNO3D GeForce RTX 3090 iCHILL X4 24 GB	€ 650 Marktplaats+1
ASUS ROG Strix RTX 3090 OC (used)	€ 800 Marktplaats+1
Some 3090 Ti listed (less common)	~€ 1000+ Marktplaats+1

What stands out

- There **are** functional 24 GB 3090 cards (various brands) going for **€650-€700** — these are among the lowest price points in recent listings and seem to frequently be “normal use, not mined, decent condition.”
 - Slightly more premium cards (partner-brand, water-cooled, or “top-end OC” versions) sit around **€700-€800**.
 - As with any second-hand market: condition, seller credibility, and tested functionality vary — some list explicitly as “works fine,” others give little info (which raises risk).
-

⌚ What a “Good Price” for 3090 Means — Based on This Market

- **€650 – €700** — this is roughly the “sweet-spot” for used 3090: reasonable VRAM headroom, and price that makes the upgrade worthwhile considering your usage (AI / image generation).
- **€700 – €800** — still fair, especially if the card is in good condition (fans working, no mining history, some minimal wear) or a nice partner-board model (better cooling, custom PCB, etc.).
- **> €850** — only worth it if it’s basically “like new,” comes with original box/accessories, maybe a short warranty or seller proof, or water-cooled / special edition — otherwise you’re not getting much extra compared to lower-priced ones.

Given what you already have (12 GB on 3060, dual-GPU 1080 Ti rig), I’d say anything \leq **€700-€750** is a *good value* for 24 GB VRAM headroom.

⚠ What to Watch Out For / When to Be Careful

- Some cards may have been used for mining — that wears down VRAM chips and fans; always ask for usage history and test under load (temps, artifacts).
- Avoid listings with little info or vague “works fine” statements — ideally prefer ones where the seller allows inspection or live testing (fan noise, temps, benchmark).
- Beware of scams: always prefer local / pick-up if possible or secure payment methods; second-hand GPUs are medium-risk compared to brand-new hardware.

- Check that card fits your system (case, PSU, connectors) — 3090 is large, draws significant power; good VRM cooling & PSU needed.
-

✓ My Recommendation (Given Your Context)

If I were you and browsing Marktplaats (or similar NL/EU second-hand markets) for a 3090:

- I'd aim for **€650–€700 first** — filter listings carefully, ask questions about usage/condition.
 - If none show up at that range, consider **€700–€750** as the upper “safe zone” before diminishing returns — especially for well-maintained partner cards.
 - I'd skip anything > €850 unless there is a very good reason (like brand-new, warranty, special cooling, water-block, etc.).
-

If you like — I can fetch **10–15 live 24 GB+ GPU listings** from *all* major NL/EU marketplaces (Marktplaats, second-hand forums, even wider EU) **translated to €** and mark which ones are “good value now,” “watch closely,” or “overpriced.”

Want me to build that list for you now?

You said:

Yes

ChatGPT said:

Cool — I dug through a number of recent resale / second-hand listings across European marketplaces (Belgium, NL, France, pan-EU sales) and found a **snapshot** of ~15–20 actual NVIDIA GeForce RTX 3090 (and similar 24 GB-VRAM 3090 boards) — prices, conditions, and what seems like “good value vs overpriced.” This should give you a better feel for what’s “reasonable now.”

🔍 Recent 3090 Listings & What They Tell Us

NVIDIA GeForce RTX 3090

€914.99

-

Rakuten - Informatique - Occasion

NVIDIA GeForce RTX 3090 (used low-cost listing)

€550.00

•

Leboncoin

NVIDIA GeForce RTX 3090 Founders Edition – Good Condition

€850.00

•

Leboncoin

NVIDIA GeForce RTX 3090 FE

€350.00

•

Leboncoin

Gainward GeForce RTX 3090 Phoenix GS 24 GB

€1,000.00

•

Leboncoin

NVIDIA GeForce RTX 3090 GeForce Vulcan OC

€1,114.99

•

Rakuten - Informatique - Occasion

Palit GeForce RTX 3090 GamingPro (LHR)

€1,014.99

•

Rakuten - Informatique - Occasion

ASUS RTX3090-24G-EK Graphics Card

€1,295.00

•

StockX

NVIDIA GeForce RTX 3090 (Rakuten – Occasion)

€914.99

•

Rakuten - Informatique - Occasion

NVIDIA GeForce RTX 3090 Founders Edition (low-price listing)

€600.00

•

Leboncoin

What stands out from the list

- The “average used-market” price for a working 24 GB 3090 tends to cluster around **~€650 – €750** — e.g. one listing at about €686 (per a recent EU price tracker) for functional used cards. [Best Value GPU](#)
 - Some partner or “good condition / lightly used” 3090 boards go for **€700-€800**. One advertisement in NL asked ~€740 for a “good-state” 3090. [Tweakers+1](#)
 - There are rarer lower-price outliers: one listing I saw was **as low as ~€550-€600** — but such low prices often demand extra caution (test thoroughly, ask for history, check for mining use or wear). [2ememain+1](#)
 - On the upper end, “like new” or “top-tier custom-PCB / partner” 3090s sometimes fetch **€800-€900+** — often because of perceived lower wear, good cooling or extras.
[leboncoin+2leboncoin+2](#)
-

✓ What I'd Call “Good Value / Smart Buy” vs “Overpriced” (Right Now)

Price range / Condition	Assessment
€600 – €650 (working card, seller lets you test)	✓ Very good value — ideal if you’re ok with potential slight wear and verify functional state before purchase.
€650 – €750 (normal used, decent condition)	✓ “Sweet-spot” — good balance of VRAM headroom + value; great for photoreal models + long-term use.
€750 – €800 (partner / custom-PCB / well-kept)	✓ Acceptable if card condition seems good and seller is reliable — gives longer lifespan potential.
€800 – €900+ (near-new, minimal wear, maybe warranty or extras)	⚠ Acceptable but less “value-efficient.” Worth it only if you need guaranteed reliability, warranty or special build; otherwise diminishing returns vs cheaper ones.
< €600 or “suspiciously cheap” / no seller history	⚠ Higher risk — only consider if you can test it thoroughly, ideally in-person before paying.

⌚ My “What I’d Do If I Were You” (Given Your Goals & Location)

If I were you and browsing right now for a 24 GB+ card for photorealistic AI generation:

- I’d aim for a 3090 priced **€650–€750**, preferably from a seller that allows **testing and verifying the card works under load** (fans, temps, driver install, memory test).
- I’d treat **€600–€650** offers as “good but potential risk” — only go for them if seller seems honest and conditions allow testing (or meet in person).
- I’d avoid anything **> €850–€900** unless it’s certified “like new” or comes from a reputable reseller — for my use case, not worth paying a big premium.

Given what you already have (12 GB cards), getting a 24 GB 3090 in the “sweet-spot” range pretty much unlocks headroom for big photoreal models and future-proofed workloads — and seems worth it *only* if price stays reasonable.

If you want — I can also **filter the list** further to **cards sold from EU sellers** only (avoid import complexity/shipping); that might be more relevant for you in Paris/Europe.

ChatGPT can make mistakes. Check important info.