

Looking at current state of ComfyUI setup on Original Win11

Current Model Distribution

From your screenshots, you have:

- CLIP/text encoders (clip folder, text_encoders)
- ControlNet models (controlnet folder with SDXL folder)
- VAE models (vae, vae_approx folders)
- Upscale models (upscale_models)
- UNet models (unet folder)
- Diffusion models (diffusion_models folder)
- Face restoration (facerestore_models)

Recommended Setup for Your Hardware

ComfyUI Server (2x GTX 1080 Ti = 22GB VRAM): This is your powerhouse for generation. Keep here:

- All diffusion models (flux1-schnell, flux1-dev, z_image_turbo)
- UNet models (flux1-dev-fp8.safetensors)
- ControlNet models (especially the large SDXL ones)
- VAE models
- Main CLIP/text encoders

Workstation (RTX 3060 12GB + Ollama): Since Ollama is already using this, keep this focused on:

- Lightweight inference support if needed
- Face restoration models (GFPGAN)
- Upscale models (RealESRGAN)
- Can serve as backup for smaller models

Optimization Tips

1. **Consolidate duplicates:** You have multiple encoder/decoder variants - keep only what you actively use
2. **Use FP8 quantized models:** I see you already have some (t5xxl_fp8_e4m3fn.safetensors) - these save significant VRAM
3. **Network access:** Ensure ComfyUI on your server is accessible from your workstation via network
4. **Model loading:** With 22GB VRAM on the server, you can load larger Flux models comfortably

Understanding Your ComfyUI Model Files

Let me break down each folder and file type you showed me:

1. CLIP Folder (Image 2)

What they are:

- `clip_1.safetensors` (240 MB) - CLIP text encoder for understanding prompts
- `t5xxl_fp8_e4m3fn.safetensors` (4.7 GB) - T5 XXL text encoder (Flux models)

What they do: Convert your text prompts into embeddings that the AI understands. Flux models need both CLIP and T5 encoders.

Can you eliminate? NO - these are essential for Flux to work. Keep both.

2. ControlNet Folder (Image 3)

What they are:

- `flux-controlnet-canny-v3.safetensors` (1.4 GB) - Edge detection control
- `flux-controlnet-depth.safetensors` (3.1 GB) - Depth map control
- **SDXL subfolder** - ControlNet models for SDXL (different model family)

What they do: Guide image generation using edge maps, depth maps, poses, etc. They give you precise control over composition.

Can you eliminate?

- Keep if you use ControlNet features (highly recommended)
 - The SDXL folder: only if you use SDXL models (not Flux)
-

3. VAE_approx Folder (Image 4)

What they are: Multiple TAESD (Tiny AutoEncoder) variants

What they do: Create quick, low-quality previews during generation so you can see progress without loading the full VAE.

Can you eliminate?

- Keep ONE that matches your model: `taesd_decoder.safetensors` for Flux
 - Delete the rest (`taefl`, `taesd3`, `taesdxl` variants unless you use those specific models)
-

4. VAE Folder (Image 5)

What they are:

- `ae.safetensors` (327 MB) - Flux AutoEncoder
- `ae-from-zimage-youtube.safetensors` (327 MB) - Duplicate/variant

What they do: Encode images to latent space and decode them back. Essential for image generation.

Can you eliminate? Keep `ae.safetensors` only. Delete the YouTube variant (it's redundant).

5. Upscale_models Folder (Image 6)

What they are:

- `Realesrgan_x4.pth` (65 MB) - 4x upscaler

What they do: Increase image resolution after generation.

Can you eliminate? Keep if you upscale images. These are small, so no harm keeping them.

6. UNet Folder (Image 7)

What they are:

- `flux1-dev-fp8.safetensors` (11.6 GB) - Flux Dev model (quantized)

What they do: This IS the actual Flux model. The core AI that generates images.

Can you eliminate? NO - this is your main model. Keep it.

7. Text_encoders Folder (Image 8)

What they are:

- `clip_1.safetensors` (240 MB) - CLIP encoder (duplicate from clip folder)
- `qwen_3_4b.safetensors` (7.8 GB) - Alternative text encoder
- `t5xxl_fp8_e4m3fn_scaled.safetensors` (5 GB) - T5 encoder variant

Can you eliminate?

- This is REDUNDANT with your clip folder
 - Keep the clip folder versions
 - Delete `qwen_3_4b` unless you specifically use it
 - Delete the "scaled" T5 variant
-

8. Facerestore_models Folder (Image 9)

What they are:

- `GFGAN-1.3.8.zip` (5.8 MB compressed)

What they do: Restore/enhance faces in generated images.

Can you eliminate? Keep if you generate faces and want enhancement. Small file, no harm keeping.

9. Diffusion_models Folder (Image 10)

What they are:

- `z_image_turbo_bf16.safetensors` (12 GB) - Fast generation model
- `flux1-schnell-fp8-e4m3fn.safetensors` (11.6 GB) - Flux Schnell (fast variant)
- `flux1-dev-kontext_fp8_scaled.safetensors` (11.6 GB) - Flux Dev variant

What they do: Different Flux model variants for different purposes (speed vs quality).

Can you eliminate?

- Keep ONE Flux model you prefer (Schnell for speed, Dev for quality)
- You already have flux1-dev in unet folder

- Delete duplicates/variants you don't use
-

10. ControlNet/SDXL Folder (Image 11)

What they are:

- controlnet-union-sdxl-1.0 subfolder

What they do: ControlNet models specifically for SDXL (not Flux).

Can you eliminate? Only if you DON'T use SDXL models. If you only use Flux, delete this entire folder.

Cleanup Recommendations

HIGH PRIORITY - Delete These:

1. Duplicate text encoders in `text_encoders` folder
2. Extra TAESD variants in `vae_approx` (keep only `taesd_decoder`)
3. `ae-from-zimage-youtube.safetensors` in VAE folder
4. Duplicate Flux models (you have `flux1-dev` in multiple places)

MEDIUM PRIORITY - Consider Deleting:

1. SDXL ControlNet folder (if you only use Flux)
2. `qwen_3_4b.safetensors` (unless you specifically use it)
3. Extra Flux variants (`Schnell`, `z_image_turbo`) if you only use one

KEEP:

- `flux1-dev-fp8.safetensors` (your main model)
 - `clip_1.safetensors` and `t5xxl_fp8_e4m3fn.safetensors` (text encoders)
 - `ae.safetensors` (VAE)
 - One TAESD decoder for previews
 - ControlNet models if you use them
 - Upscale models
-

File Sharing & Workflow Access Plan

Problem:

Your MacBook Air M2 can access ComfyUI server via web, but:

- Saved workflows → Go to server's local drive ✓
- Loading workflows → Must be on Mac's local drive X
- Can't access workflows saved on the server

Solution Options:

Option 1: OneDrive Sync (Recommended for You)

Since you already use OneDrive:

On ComfyUI Server:

1. Move workflow folder to OneDrive
2. Create symlink in ComfyUI directory:

```
mklink /D "D:\Misc\ComfyUI\ComfyUI\user\default\workflows"  
"C:\Users\YourName\OneDrive\ComfyUI_Workflows"
```

On MacBook Air M2:

1. Install OneDrive
2. Let workflows sync automatically
3. Point ComfyUI to OneDrive folder

Pros:

- Automatic sync across all devices
- Works over internet
- No VPN needed
- Already familiar with OneDrive

Cons:

- Uses cloud storage quota
 - Requires internet for sync
-

My Recommendation:

Use OneDrive sync since you're already using it. Here's the setup:

Step-by-Step Setup:

1. On ComfyUI Server:

1. Create folder: C:\Users\YourName\OneDrive\ComfyUI_Shared\workflows
2. Copy existing workflows there
3. Open Command Prompt as Administrator
4. Run: mklink /D "D:\Misc\ComfyUI\ComfyUI\user\default\workflows" "C:\Users\YourName\OneDrive\ComfyUI_Shared\workflows"

2. On MacBook Air M2:

1. Install OneDrive (if not already)
2. Let ComfyUI_Shared folder sync
3. In ComfyUI settings, point workflow directory to:
~/Library/CloudStorage/OneDrive/ComfyUI_Shared/workflows

3. On Windows Workstation (optional):

Same as MacBook - let OneDrive sync and point to that folder.

For Model Sharing Between Windows Machines:**

****Option A: Network Path****

On ComfyUI server, share models folder:

Share: D:\Misc\ComfyUI\ComfyUI\models

Access from workstation: \\server-ip\models

In ComfyUI extra_model_paths.yaml:

yaml

workstation:

```
base_path: \\server-ip\models\  
checkpoints: diffusion_models  
vae: vae  
clip: clip  
controlnet: controlnet
```

Option B: Duplicate Essential Models Keep frequently used models on both machines, sync

Best Method: Upload Directly Here

1. Zip your workflows folder:

powershell

```
# Navigate to your workflows directory  
cd "D:\Misc\ComfyUI\ComfyUI\user\default\workflows"  
  
# Create a zip file  
Compress-Archive -Path *.json -DestinationPath "C:\Users\YourName\Desktop\my_workflows.zip"
```

2. Upload the zip here - I can analyze all your workflow JSON files to see which models they reference

What I'll Check For:

When you send the workflows, I'll scan for:

- **Checkpoint/model names** - which Flux variants you actually use
- **ControlNet references** - if you use SDXL ControlNets
- **VAE references** - which VAE/TAESD models are needed
- **Text encoder requirements** - if you use alternative encoders like Qwen
- **Custom nodes** - that might require specific models

Quick Manual Check You Can Do:

Open any workflow JSON file in a text editor and search for these terms:

- "qwen" - if found, keep qwen_3_4b.safetensors
- "sdxl" - if found, keep SDXL ControlNet folder
- "schnell" - if found, keep flux1-schnell model
- "z_image_turbo" - if found, keep that model
- "ae-from-zimage" - if found, keep that VAE variant

Models Actually Used in Your Workflows:

CLIP/Text Encoders:

- **clip_l.safetensors** - Used in most Flux workflows
- **t5xxl_fp8_e4m3fn.safetensors** - Used in Flux workflows
- **qwen_3_4b.safetensors** - NOT found in any workflow
- **t5xxl_fp8_e4m3fn_scaled.safetensors** - NOT found in any workflow

Main Models:

- **flux1-dev-fp8.safetensors** (or flux1-dev) - Used in multiple workflows
- **flux1-schnell** - Used in several "schnell" workflows
- **z_image_turbo** - Only found in "zimage.json" workflow
- **flux1-dev-kontext** - NOT found in any workflow

VAE:

- **ae.safetensors** - Used in Flux workflows
- **ae-from-zimage-youtube** - Only in zimage.json workflow

SDXL Models:

- **SDXL ControlNet folder** - Found in "Simple SDXL Workflow" files, but you seem focused on Flux

- Keep if you use those "Simple SDXL Workflow" files

ControlNet:

- **flux-controlnet-canny-v3** and **flux-controlnet-depth** - Used in some workflows

Safe to Delete:

Based on your actual workflows:

1. ~~X~~ **qwen_3_4b.safetensors** (7.8 GB) - Not used
2. ~~X~~ **t5xxl_fp8_e4m3fn_scaled.safetensors** (5 GB) - Not used
3. ~~X~~ **flux1-dev-kontext_fp8_scaled.safetensors** (11.6 GB) - Not used
4. ~~X~~ **Extra TAESD variants** in vae_approx (except taesd_decoder)
5. ~~X~~ **ae-from-zimage-youtube.safetensors** - Unless you use zimage.json

Keep if you use specific workflows:

- **z_image_turbo_bf16.safetensors** - Only if you use zimage.json
- **flux1-schnell-fp8** - Keep if you use the "schnell" fast workflows
- **SDXL ControlNet** - Only if you use "Simple SDXL Workflow" files

Updated Space Savings:

- Removing unused files: **~24 GB** (safer estimate)
- If you also remove flux1-schnell (keeping only flux1-dev): **~35 GB**